

Install library: `readxl`

Time Series Regression

06

Forecasting: Principles and Practice

Preface

1 Getting started

2 Time series graphics

3 Time series decomposition

4 Time series features

5 The forecaster's toolbox

6 Judgmental forecasts

7 Time series regression models

8 Exponential smoothing

9 ARIMA models

10 Dynamic regression models

Outline

- (10) Introduction
 - Time series (TS) linear regression: New Car Sales Example
- TS Forecasting with regression
 - (20) Excel
 - (20) R
 - (20) ex-ante vs ex-post forecasts
- (30) Regression assumptions: Diagnostic tests
- (15) Model selection
 - Useful predictors
- Appendix: Multiple regression-matrix notation



Multiple linear **regression** model

- Previous classes: time series forecasting. Now, regression.
- $y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \dots + \beta_k x_{ki} + \epsilon_i$
where k is the number of predictors.
 - If $k = 1$, the model is called a **simple** linear regression.
 - If $k \geq 2$, the model is called a **multiple** linear regression.
- y_i **dependent** variable; **response** variable
- x_{ki} the i th observation on the k th **independent variable**
 - x_k is also called the **predictor** variable or the **explanatory variable**.
- Errors $\epsilon_i \sim NID(0, \sigma^2)$



Time series regression Example

- **Dependent** variable
(ML terms: **target** variable)
 - New Car Sales (NCS)
- **Independent** variables (**explanatory** variables)
(ML terms: **predictors, features**)
 - Disposable personal income per capita (DPIPC)
 - Unemployment rate (UR)
 - Bank prime rate (PR)
 - University of Michigan Index of Consumer Sentiment (UMICS)

Sheet: NewCarSales

Date	NCS	DPIPC	UR	PR	UMICS
Feb-02	152641	27840	5.7	4.75	93.13
May-02	166036	28134	5.83	4.75	94.1
Aug-02	175863	28168	5.73	4.75	87.27
Nov-02	151219	28363	5.87	4.45	83.83
...					
Feb-16	204681	42807	4.93	3.5	91.57
May-16	221073	43265	4.87	3.5	92.4
Aug-16	229123	43651	4.9	3.5	90.33
Nov-16	221392	43759	4.7	3.55	93.07


The number of rows: $n = 60$.

How many years?

Positive or negative relationship?



Example



```
d <- read_excel("09-TS-regression.xlsx", sheet = "NewCarSales")

ncs <- d %>%
  mutate(Quarter = yearquarter("2002 Q1") + 0:59) %>%
  select(-Date) %>%
  relocate(Quarter) %>%
  as_tsibble(index = Quarter)


ncs %>% autoplot(NCS)
```

Sheet: NewCarSales

Date	NCS	DPIPC	UR	PR	UMICS
Feb-02	152641	27840	5.7	4.75	93.13
May-02	166036	28134	5.83	4.75	94.1
Aug-02	175863	28168	5.73	4.75	87.27
Nov-02	151219	28363	5.87	4.45	83.83
...					
Feb-16	204681	42807	4.93	3.5	91.57
May-16	221073	43265	4.87	3.5	92.4
Aug-16	229123	43651	4.9	3.5	90.33
Nov-16	221392	43759	4.7	3.55	93.07

```
> ncs
# A tsibble: 60 x 6 [1Q]
   Quarter      NCS DPIPC   UR   PR UMICS
   <qtr>      <dbl> <dbl> <dbl> <dbl> <dbl>
1 2002 Q1  152641  27840  5.7   4.75  93.1
2 2002 Q2  166036  28134  5.83  4.75  94.1
3 2002 Q3  175863  28168  5.73  4.75  87.3
4 2002 Q4  151219  28363  5.87  4.45  83.8
5 2003 Q1  152843  28584  5.87  4.25  80.0
6 2003 Q2  173360  28958  6.13  4.24  89.3
7 2003 Q3  178086  29539  6.13  4     89.3
8 2003 Q4  154916  29708  5.83  4     92.0
9 2004 Q1  159086  30094  5.7   4     98
10 2004 Q2  172174  30537  5.6   4     93.3
# ... with 50 more rows
# i Use `print(n = ...)` to see more rows
```

Example: Time plot




NCS and UMICS
Positive or negative relationship?

```
ncs %>%  
  pivot_longer(-Quarter) %>%  
  ggplot(aes(x= Quarter, y = value, color = name)) +  
  geom_line() +  
  facet_grid(name~., scales = "free_y")      # try without facet
```

Scatterplots & correlation coefficients

```
ncs %>%  
  GGally::ggpairs(columns = 2:6)
```

NCS and UMICS
Positive or negative relationship?





Example: Regression equation

Regression equation

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 + \cdots + b_kx_k$$

where

- b_0 is the **constant term**—the value of y when the values of all the x variables are set at 0;
- $b_1, \dots, b_k$ are the **regression coefficients (slopes)**

Regression equation

$$\text{NCS.Forecast} = 25428.2554 + 3.2275 \cdot \text{DPIPC} - 7666.8516 \cdot \text{UR} - 3219.4991 \cdot \text{PR} + 1122.4290 \cdot \text{UMICS}$$

```
fit <- ncs %>%  
  model(tslm = TSLM(NCS ~ DPIPC + UR + PR + UMICS))  
report(fit)
```

Coefficients:

	Estimate
(Intercept)	25428.2554
DPIPC	3.2275
UR	-7666.8516
PR	-3219.4991
UMICS	1122.4290



	Coefficients
Intercept	25428.25537
DPIPC	3.227461566
UR	-7666.851578
PR	-3219.499125
UMICS	1122.429026



Excel
Data Analysis → Regression



Example: Regression equation

Coefficient interpretation

Regression equation

NCS.Forecast = 25428.2554 + 3.2275*DPIPC – 7666.8516*UR – 3219.4991*PR + 1122.4290*UMICS

Regression equation

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 + \cdots + b_kx_k$$

where

- b_0 is the constant term—the value of y when the values of all the x variables are set at 0
- $b_1, \dots, b_k$ are the regression coefficients (slopes), indicating the change in the value of y for each unit of change in the relevant x value *ceteris paribus* (given that all other things being equal)
 - Holding other variables (DPIPC, UR, PR) constant, a one-unit increase in UMICS (the consumer sentiment index) is associated with an increase of 1122.43 units in NCS (number of cars sold) on average.



Excel's Data Analysis: Regression

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T
1	Date	NCS	DPIPC	UR	PR	UMICS	SUMMARY OUTPUT													
2	02-Feb	152641	27840	5.7	4.75	93.13								Coefficients	25428.26	3.227462	-7666.85	-3219.5	1122.429	
3	02-May	166036	28134	5.83	4.75	94.1	Regression Statistics							Seasonal Naïve for 2017	Ones	DPIPC	UR	PR	UMICS	NCS.Forecast
4	02-Aug	175863	28168	5.73	4.75	87.27	Multiple F	0.910035						Q1	1	42807	4.93	3.5	91.57	217301.2
5	02-Nov	151219	28363	5.87	4.45	83.83	R Square	0.828163						Q2	1	43265	4.87	3.5	92.4	220171.0
6	03-Feb	152843	28584	5.87	4.25	79.97	Adjusted R	0.815666						Q3	1	43651	4.9	3.5	90.33	218863.4
7	03-May	173360	28958	6.13	4.24	89.27	Standard Error	11689.89						Series: NCS						
8	03-Aug	178086	29539	6.13	4	89.3	Observations	60						Model: TSLM						
9	03-Nov	154916	29708	5.83	4	91.97								Residuals:						
10	04-Feb	159086	30094	5.7	4	98	ANOVA							Min	1Q	Median	3Q	Max		
11	04-May	172174	30537	5.6	4	93.33		df	SS	MS	F			-27502.7	-8272.9	488.2	7831.5	20792.9		
12	04-Aug	178077	30800	5.43	4.42	95.6	Regression	4	3.62E+10	9.06E+09	66.26779									
13	04-Nov	161216	31353	5.43	4.94	93.87	Residual	55	7.52E+09	1.37E+08										
14	05-Feb	160070	31145	5.3	5.44	94.07	Total	59	4.37E+10					Coefficients:						
15	05-May	183594	31533	5.1	5.91	90.2								Estimate	Std. Error	t value	Pr(> t)			
16	05-Aug	184878	31961	4.97	6.43	87.5	Coefficients Standard Error t Stat P-value							(Intercept)	25428.2554	37465.7657	0.679	0.5002		
17	05-Nov	153405	32396	4.97	6.97	82.43	Intercept	25428.26	37465.77	0.678706	0.50017		DPIPC	3.2275	0.3916	8.242	3.52e-11	***		
18	06-Feb	164864	33217	4.73	7.43	88.93	DPIPC	3.227462	0.391571	8.242334	3.52E-11		UR	-7666.8516	1662.4019	-4.612	2.42e-05	***		
19	06-May	178363	33448	4.63	7.9	83.8	UR	-7666.85	1662.402	-4.61191	2.42E-05		PR	-3219.4991	1412.5314	-2.279	0.0266	*		
20	06-Aug	181615	33696	4.63	8.25	84.03	PR	-3219.5	1412.531	-2.27924	0.026556		UMICS	1122.4290	210.6755	5.328	1.91e-06	***		
21	06-Nov	160780	33991	4.43	8.25	92.47	UMICS	1122.429	210.6755	5.327763	1.91E-06		---							
22	07-Feb	167528	34457	4.5	8.25	92.2								Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1						
23	07-May	180006	34720	4.5	8.25	86.9								Residual standard error: 11690 on 55 degrees of freedom						
														Multiple R-squared: 0.8282, Adjusted R-squared: 0.8157						
														F-statistic: 66.27 on 4 and 55 DF, p-value: < 2.22e-16						



Forecasting with regression

Scenario-based forecast

ex-ante vs ex-post forecasts



TS Forecasting with Regression

Date	NCS	DPIPC	UR	PR	UMICS
Feb-02	152641	27840	5.7	4.75	93.13
May-02	166036	28134	5.83	4.75	94.1
Aug-02	175863	28168	5.73	4.75	87.27
Nov-02	151219	28363	5.87	4.45	83.83
...					
Feb-16	204681	42807	4.93	3.5	91.57
May-16	221073	43265	4.87	3.5	92.4
Aug-16	229123	43651	4.9	3.5	90.33
Nov-16	221392	43759	4.7	3.55	93.07

	<i>Coefficients</i>
Intercept	25428.25537
DPIPC	3.227461566
UR	-7666.851578
PR	-3219.499125
UMICS	1122.429026



- Regression equation

$$\text{NCS.Forecast} = 25428.2554 + 3.2275 \cdot \text{DPIPC} - 7666.8516 \cdot \text{UR} - 3219.4991 \cdot \text{PR} + 1122.4290 \cdot \text{UMICS}$$

- Suppose that in 2017 Q1, the forecasts of independent variables (i.e., DPIPC, UR, PR, UMICS) come from *seasonal naïve*.

$$\text{DPIPC} = 42807,$$

$$\text{UR} = 4.93,$$

$$\text{PR} = 3.50,$$

$$\text{UMICS} = 91.57.$$

- Different TS forecasting methods can be used for different series.
- Then, the NCS forecast for 2017Q1 is

$$\begin{aligned} \text{NCS.Forecast} &= 25428.2554 + 3.2275(42807) - 7666.8516(4.93) - 3219.4991(3.50) + 1122.4290(91.57) \\ &= 217301 \end{aligned}$$

Forecast 2017 Q2, Q3, Q4

Forecasting with Regression

Forecasts of predictors

- We require the forecasts of the predictors.
- To obtain these forecasts,
 - We can use one of the simple methods or more sophisticated pure time series approaches.
 - Alternatively, forecasts from other sources may be available and can be used.
 - Forecasts from government agency
 - GDP
 - Weather
 - Forecasts from sales department in the Sales and Operations Planning (S&OP) meeting.



R: TS cross validation to select forecast for predictors*



```
# "unpivot": long format
ncs_long <- ncs %>%
  select(-NCS) %>%
  pivot_longer(!Quarter, names_to="variable", values_to = "value")

ncs_long

trCV <- ncs_long %>%
  slice(1:(n()-4*4)) %>%
  stretch_tsibble(.init=6, .step=1) %>%
  relocate(.id, Quarter)

trCV %>%
  model(
    Mean = MEAN(value),
    Naive = NAIVE(value),
    Drift = NAIVE(value ~ drift()),
    Snaive = SNAIVE(value)
  ) %>%
  forecast(h=4) %>%
  accuracy(ncs_long) %>%
  arrange(variable, RMSE)
```

!Quarter All data except column Quarter

```
# A tsibble: 240 x 3 [1Q]
# Key:   variable [4]
  Quarter variable value
  <ctr>   <chr>   <dbl>
1 2002 Q1 DPIPc    27840
2 2002 Q1 UR       5.7
3 2002 Q1 PR       4.75
4 2002 Q1 UMICS    93.1
5 2002 Q2 DPIPc    28134
6 2002 Q2 UR       5.83
7 2002 Q2 PR       4.75
8 2002 Q2 UMICS    94.1
9 2002 Q3 DPIPc    28168
10 2002 Q3 UR       5.73
# ... with 230 more rows
```

	.model	variable	.type	ME	RMSE	MAE	MPE	MAPE	MASE	RMSSE	ACF1	Predictor
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	
1	Drift	DPIPc	Test	-28.7	634.	430.	-0.0849	1.17	0.360	0.484	0.664	DPIPc
2	Naive	DPIPc	Test	677.	955.	814.	1.85	2.22	0.680	0.729	0.485	
3	Snaive	DPIPc	Test	1088.	1324.	1213.	3.00	3.33	1.01	1.01	0.605	
4	Mean	DPIPc	Test	4649.	4950.	4649.	12.3	12.3	3.88	3.78	0.915	
5	Naive	PR	Test	-0.0290	0.893	0.471	-2.19	9.27	0.616	0.708	0.792	PR
6	Drift	PR	Test	-0.0409	0.977	0.560	-1.34	11.6	0.733	0.775	0.779	
7	Snaive	PR	Test	-0.0622	1.27	0.755	-4.53	15.1	0.988	1.00	0.915	
8	Mean	PR	Test	-0.543	1.92	1.76	-26.4	42.8	2.30	1.52	0.967	
9	Naive	UMICS	Test	0.0504	7.61	5.70	-0.486	7.57	0.727	0.757	0.678	UMIC
10	Drift	UMICS	Test	1.11	8.10	6.25	0.978	8.21	0.797	0.807	0.682	
11	Snaive	UMICS	Test	0.319	9.85	7.56	-0.528	10.0	0.964	0.981	0.669	
12	Mean	UMICS	Test	-4.24	12.0	9.39	-7.34	12.9	1.20	1.20	0.876	
13	Naive	UR	Test	-0.0561	0.873	0.579	-1.77	8.43	0.658	0.712	0.796	UR
14	Drift	UR	Test	-0.113	0.926	0.644	-2.18	9.26	0.732	0.756	0.786	
15	Snaive	UR	Test	-0.0780	1.24	0.904	-2.92	13.0	1.03	1.01	0.909	
16	Mean	UR	Test	0.575	1.93	1.50	2.44	21.0	1.71	1.58	0.966	

Based on
smallest
RMSE

FC

Drift

Naive

Naive

Naive

Also try ARIMA and ETS

“What-If” Scenarios

- Suppose the sales of a company depend on the price it charges and the amount it spends on different media:

$$\text{Sales} = 120 - 0.4(\text{Avg Price}) + 5(\text{Media Spending})$$

where sales and media spending are given in millions of dollars

- Current plan: Avg price = 10 USD/unit, Media spending = 2 million.
 - Sales.Forecast = $120 - 0.4(10) + 5(2) = 126$ million
- But, management wants to hit the sales target of 130 million.
- If we keep avg price and change media spending, how much must be spent on the media to hit the target?

$$130 = 120 - 0.4(10) + 5(x_2) \rightarrow x_2 = 2.8 \text{ million}$$

Ex-ante vs Ex-post forecasts

Ex-ante uses forecasted predictors, Ex-post uses actual predictors


Ex-ante forecast

Ex-ante forecasts are those that are made using only the information that is available.

- Based on forecasts rather than actual results.

	Intercept	DPIPC	UR	PR	UMICS	NCS.FC
	25428.26	3.227462	-7666.85	-3219.5	1122.429	
2017Q1		42807	4.93	3.5	91.57	217,301
2017Q2		Forecasts using data up to 2016				220,171
2017Q3		43051	4.5	3.3	90.55	218,863
2017Q4		43759	4.7	3.55	93.07	223,660

$$\text{NCS.FC} = 25428.26 + 3.23 * \text{DPIPC.FC} - 7666.85 * \text{UR.FC} - 3219.5 * \text{PR.FC} + 1122.43 * \text{UMICS.FC}$$



Ex-post forecast

Ex-post forecasts are those that are made using later information on the predictors.

- Based on actual results rather than forecasts

	Intercept	DPIPC	UR	PR	UMICS	NCS.FS
	25428.26	3.227462	-7666.85	-3219.5	1122.429	
2017Q1		42807	4.97	3.4	92.73	218,618
2017Q2		43211	4.8	3.4	92.73	218,314
2017Q3		43651	4.8	3.4	92.73	223,220
2017Q4		43759	4.6	3.4	5.1	126,169

Actual values in 2017

$$\text{NCS.FC} = 25428.26 + 3.23 * \text{DPIPC.Actual} - 7666.85 * \text{UR.Actual} - 3219.0 * \text{PR.Actual} + 1122.43 * \text{UMICS.Actual}$$


Ex-ante vs Ex-post forecasts

- A comparative evaluation of ex-ante forecasts and ex-post forecasts can help to separate out the sources of forecast uncertainty.
- This will show whether forecast errors have arisen
 - a) due to poor forecasts of the predictor or
 - b) [*model misspecification*] due to a poor forecasting model (regression model, here).

Ex-ante forecast

Ex-post forecast

2016 Q4
Occurrence of Event



	Intercept	DPIPC	UR	PR	UMICS	NCS.FC
	25428.26	3.227462	-7666.85	-3219.5	1122.429	
2017Q1		42807	4.93	3.5	91.57	217,301
2017Q2		Forecasts using data up to 2016				220,171
2017Q3		43051	4.5	3.3	90.55	218,863
2017Q4		43759	4.7	3.55	93.07	223,660

	Intercept	DPIPC	UR	PR	UMICS	NCS.FS
	25428.26	3.227462	-7666.85	-3219.5	1122.429	
2017Q1		42807	4.97	3.4	92.73	218,618
2017Q2		43201	4.9	3.4	92.73	218,314
2017Q3		43651	4.6	3.4	92.73	223,220
2017Q4		43759	4.6	3.4	5.1	126,169

Actual values in 2017

$$\text{NCS.FC} = 25428.26 + 3.23 * \text{DPIPC.FC} - 7666.85 * \text{UR.FC} - 3219.5 * \text{PR.FC} + 1122.43 * \text{UMICS.FC}$$

$$\text{NCS.FC} = 25428.26 + 3.23 * \text{DPIPC.Actual} - 7666.85 * \text{UR.Actual} - 3219.0 * \text{PR.Actual} + 1122.43 * \text{UMICS.Actual}$$

Example

- I) due to poor forecasts of the predictor: MAPE ex-ante forecast 60%,
- II) due to a poor forecasting model: MAPE ex-ante forecast 60%,

while MAPE ex-post forecast 8%
while MAPE ex-post forecast 58%

Ex-ante forecast: Prediction interval*

- The prediction intervals for ex-ante forecasts do not include the uncertainty associated with the future values of the predictor variables.
- They assume that the values of the predictors are known in advance.

```
ncs_tr <- ncs %>%
  filter(year(Quarter) <= 2015)


ncs_test <- ncs %>%
  filter(year(Quarter) >= 2016)

fit <- ncs_tr %>%
  model(tslm = TSLM(NCS ~ DPIPc + UR + PR + UMICS))

fc_predictor <- tsibble(
  Quarter = yearquarter("2016 Q1") + 0:3,
  DPIPc = rep(42621, 4),
  PR = rep( 3.29, 4),
  UMICS = rep(91.3, 4),
  UR = rep(5, 4),
  index = Quarter
)

fc_exante <- forecast(fit, new_data = fc_predictor)

fc_exante %>%
  hilo(level=99)
```



.model	Quarter	NCS	.mean	DPIPc	PR	UMICS	UR	~99%~
<chr>	<qtr>	<dist>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<hilo>
1 tslm	2016 Q1	N(217037, 1.6e+08)	217037.	42621	3.29	91.3	5	[184181.4, 249893]99
2 tslm	2016 Q2	N(217037, 1.6e+08)	217037.	42621	3.29	91.3	5	[184181.4, 249893]99
3 tslm	2016 Q3	N(217037, 1.6e+08)	217037.	42621	3.29	91.3	5	[184181.4, 249893]99
4 tslm	2016 Q4	N(217037, 1.6e+08)	217037.	42621	3.29	91.3	5	[184181.4, 249893]99

Regression model: Assumptions



Multiple Linear Regression: Assumptions

$$y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \cdots + \beta_k x_{ki} + \epsilon_i$$

- The relationship between y and x_k is linear and given above.
- The x 's are nonstochastic variables. In addition, no exact linear relationship exists between two or more independent variables.
- Errors $\epsilon_i \sim NID(0, \sigma_e^2)$
 - The error has zero expected value for all observations
 - **Homoscedasticity**: Error term has constant variance for all observations.
 - **White noise**: Errors corresponding to different observations are independent and therefore uncorrelated. No serial correlation.
 - The error term is normally distributed.



Diagnostic Test

1. Do signs of coefficients make sense?
2. Is the coefficient statistically significant?
3. Explanatory power: **Adjusted R-Square**
4. Residual analysis
5. Multicollinearity

Series: NCS
Model: TSLM

Residuals:
Min 1Q Median 3Q Max
-27502.7 -8272.9 488.2 7831.5 20792.9

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 25428.2554 37465.7657 0.679 0.5002
DPIPC 3.2275 0.3916 8.242 3.52e-11 ***
UR -7666.8516 1662.4019 -4.612 2.42e-05 ***
PR -3219.4991 1412.5314 -2.279 0.0266 *
UMICS 1122.4290 210.6755 5.328 1.91e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 11690 on 55 degrees of freedom
Multiple R-squared: 0.8282, Adjusted R-squared: 0.8157
F-statistic: 66.27 on 4 and 55 DF, p-value: < 2.22e-16

Regression Statistics								
Multiple R	0.910034731							
R Square	0.828163212							
Adjusted R	0.815665991							
Standard E	11689.89219							
Observations	60							
ANOVA								
	df	SS	MS	F	Significance F			
Regression	4	3.62E+10	9.06E+09	66.26779	2.2E-20			
Residual	55	7.52E+09	1.37E+08					
Total	59	4.37E+10						
Coefficients								
	Estimate	Standard Error	t Stat	P-value	Lower 95%	Upper 95%	Lower 95.0%	Upper 95.0%
Intercept	25428.25537	37465.77	0.678706	0.50017	-49654.8	100511.3	-49654.8	100511.3
DPIPC	3.227461566	0.391571	8.242334	3.52E-11	2.442735	4.012188	2.442735	4.012188
UR	-7666.851578	1662.402	-4.61191	2.42E-05	-10998.4	-4335.32	-10998.4	-4335.32
PR	-3219.499125	1412.531	-2.27924	0.026556	-6050.28	-388.723	-6050.28	-388.723
UMICS	1122.429026	210.6755	5.327763	1.91E-06	700.226	1544.632	700.226	1544.632

1. Logic: Check the signs of the coefficients

$$\text{Predicted Sale} = 25304.35 + 3.227 \cdot \text{DPIPC} - 7659.605 \cdot \text{UR} - 3216.77 \cdot \text{PR} + 1123.363 \cdot \text{UMICS}$$

Dependent variable

- NCS = New car sales

Independent variables

- DPIPC = Disposable income per capita
- UMICS = University of Michigan Index of Consumer Sentiment
- UR = Unemployment rate
- PR = Prime interest rate

These coefficients measure the marginal effects, i.e., the effect of each predictor given all other things being equal (*ceteris paribus*).

2. Check the significance of the coefficients

Example: NCS

- Null hypothesis $H_0: \beta_i = 0$
- Alternative hypothesis $H_1: \beta_i \neq 0$

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	25304.3496	37454.5300	0.676	0.5021
DIPIC	3.2273	0.3916	8.240	3.54e-11 ***
UR	-7659.6055	1661.5855	-4.610	2.44e-05 ***
PR	-3216.7766	1412.4494	-2.277	0.0267 *
UMICS	1123.3628	210.5820	5.335	1.86e-06 ***

signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

- The ***t* statistic** is the ratio of an estimated coefficient to its standard error.
- For DIPIC, $\hat{\beta}_1 = 3.2273$, and its standard error is 0.3916, the *t* value = $3.2273/0.3916 = 8.24$



2. Check the significance of the coefficients

Example: NCS

- Null hypothesis $H_0: \beta_i = 0$
- Alternative hypothesis $H_1: \beta_i \neq 0$

- The **t statistic** follows the t distribution with the degrees of freedom, T_{df}
 $df = n - (k + 1) = 60 - (4 + 1) = 55$.

- For PR, $t = -2.277$, the p-value is
 $\Pr(|T_{55}| > 2.277) = 2(0.01335) = 0.0267$


```
> 1-pt(2.277, 55) # right tail
[1] 0.01334926
> 2*(1-pt(2.277, 55)) # two tail
[1] 0.02669851
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	25304.3496	37454.5300	0.676	0.5021	
DPIPC	3.2273	0.3916	8.240	3.54e-11	***
UR	-7659.6055	1661.5855	-4.610	2.44e-05	***
PR	-3216.7766	1412.4494	-2.277	0.0267	*
UMICS	1123.3628	210.5820	5.335	1.86e-06	***

 signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Three stars (or asterisks) represent a highly significant p-value.



$$H_0: \beta_4 = 0 \quad \text{vs} \quad H_1: \beta_4 \neq 0$$

Let significance level $\alpha = 0.05$. Since p-value = $1.86 \times 10^{-6} < \alpha$, there is sufficient evidence to reject the null hypothesis H_0 in favor of H_1 .
 Remove PR: Re-do the regression.



3. Coefficient of determination R^2

TABLE 6 ANOVA Table for Multiple Regression

Source	Sum of Squares	df	Mean Square	F Ratio
Regression	SSR	k	$MSR = SSR/k$	$F = \frac{MSR}{MSE}$
Error	SSE	$n - k - 1$	$MSE = SSE/(n - k - 1)$	
Total	SST	$n - 1$		

Coefficients:

```

      Estimate Std. Error t value Pr(>|t|)
(Intercept) 25304.3496 37454.5300  0.676  0.5021
DPIPC        3.2273    0.3916   8.240 3.54e-11 ***
UR        -7659.6055  1661.5855  -4.610 2.44e-05 ***
PR        -3216.7766  1412.4494  -2.277  0.0267 *
UMICS       1123.3628   210.5820   5.335 1.86e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 11690 on 55 degrees of freedom

Multiple R-squared: 0.8281, Adjusted R-squared: 0.8156

F-statistic: 66.26 on 4 and 55 DF, p-value: < 2.2e-16

The model explains 82.81% of the variation in NCS.

Error sum of squares $SSE = \sum (y_i - \hat{y}_i)^2$
 Total sum of squares $SST = \sum (y_i - \bar{y})^2$

$$R^2 = \frac{\text{Explained variation in Y}}{\text{Total variation in Y}} = \frac{SSR}{SST}$$

$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2} = 1 - \frac{SSE}{SST}$$

Regression Statistics	
Multiple R	0.910034731
R Square	0.828163212
Adjusted R Square	0.815665991
Standard Error	11689.89219
Observations	60

Adjusted R^2 : $\bar{R}^2$

- Adding any variable increases the value of R^2 , even if that variable is irrelevant.
- For these reasons, forecasters should not use R^2 to determine whether a model will give good predictions, as it will lead to **overfitting**.
- An alternative which is designed to overcome these problems is the **adjusted R^2** .

$$\bar{R}^2 = 1 - (1 - R^2) \frac{n - 1}{n - k - 1}$$

where n is the number of observations, and k is the number of predictors.

- Adjusted R^2 no longer increases with each added predictor.
- Using this measure, the best model will be the one with the largest value of adjusted R^2 .



Adjusted R^2

Underfitting vs overfitting

```
> fitUnder <- ncs %>%
+   model(ts1m = TSLM(NCS ~ DPIPC))
>
> report(fitUnder)
Series: NCS
Model: TSLM
```

```
Residuals:
    Min      1Q  Median      3Q      Max
-56557 -15584   3961  18723  37605
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  6.640e+04  2.505e+04   2.651 0.010329 *
DPIPC        2.866e+00  6.917e-01   4.144 0.000112 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 24120 on 58 degrees of freedom
Multiple R-squared: 0.2284,    Adjusted R-squared: 0.2151
F-statistic: 17.17 on 1 and 58 DF, p-value: 0.00011246
```

```
> fit <- ncs %>%
+   model(ts1m = TSLM(NCS ~ DPIPC + UR + PR + UMICS))
>
> report(fit)
Series: NCS
Model: TSLM
```

```
Residuals:
    Min      1Q  Median      3Q      Max
-27502.7 -8272.9   488.2  7831.5  20792.9
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 25428.2554 37465.7657   0.679  0.5002
DPIPC         3.2275    0.3916   8.242 3.52e-11 ***
UR          -7666.8516  1662.4019  -4.612 2.42e-05 ***
PR          -3219.4991  1412.5314  -2.279  0.0266 *
UMICS        1122.4290   210.6755   5.328 1.91e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 11690 on 55 degrees of freedom
Multiple R-squared: 0.8282,    Adjusted R-squared: 0.8157
F-statistic: 66.27 on 4 and 55 DF, p-value: < 2.22e-16
```

```
> ncs <- ncs %>%
+   mutate(URPR = UR*PR)
> fitover <- ncs %>%
+   model(ts1m = TSLM(NCS ~ DPIPC + UR + PR + UMICS + URPR))
> report(fitover)
Series: NCS
Model: TSLM
```

```
Residuals:
    Min      1Q  Median      3Q      Max
-26857 -7770   506   7984  20958
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 22495.4038 38639.1676   0.582 0.562861
DPIPC         3.0201    0.6995   4.318 6.79e-05 ***
UR          -3562.6781  11552.0453  -0.308 0.758963
PR          3019.8538  17434.4371   0.173 0.863132
UMICS       1064.7210    266.3214   3.998 0.000195 ***
URPR        -1451.6013  4042.6164  -0.359 0.720939
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 11780 on 54 degrees of freedom
Multiple R-squared: 0.8286,    Adjusted R-squared: 0.8127
F-statistic: 52.2 on 5 and 54 DF, p-value: < 2.22e-16
```

Adding any variable increases the value of R^2 , even if that variable is irrelevant.

TRUE/FALSE: " R^2 tells you the % of variation explained by only the variables that *actually affect* the dependent variable"



For these reasons, forecasters should not use R^2 to determine whether a model will give good predictions, as it will lead to **overfitting**.

Adjusted R^2 no longer increases with each added predictor.

Using this measure, the best model will be the one with the largest value of adjusted R^2 .


4. Check assumptions of errors

Homoscedasticity: No changing variation over time.







Residual plots against predictors

```
ncs %>%  
  left_join(residuals(fit), by = "Quarter") %>%  
  pivot_longer(DPIPC:UMICS,  
               names_to = "regressor", values_to = "x") %>%  
  ggplot() +  
  geom_point(aes(x=x, y=.resid)) +  
  facet_wrap(. ~regressor, scales = "free_x") +  
  labs(y="Residuals", x="")
```




- Examine scatterplots of residuals against each of the predictor variables
- If these scatter plots show a pattern, then the relationship may be nonlinear and the model will need to be modified.
- It is also necessary to plot the residuals against any predictors that are not in the model.
- If any of these show a pattern, then the corresponding predictor may need to be added.



Excel

Residual plots against fitted values


```
augment(fit) %>%  
  ggplot() +  
  geom_point(aes(x=.fitted, y = .resid)) +  
  labs(x="Fitted", y="Residuals")
```



The random scatter suggests the errors are homoskedastic.



4. Check assumptions of errors: Serial correlation



```
fit %>% gg_tsresiduals()
```

```
augment(fit) %>%  
  features(.innov, ljung_box, lag=10)
```

	.model	lb_stat	lb_pvalue
	<chr>	<dbl>	<dbl>
1	ts1m	79.0	7.92e-13

Let $\rho_j = \text{corr}(\epsilon_t, \epsilon_{t-j})$ be the autocorrelation at lag j
 $H_0: \rho_1 = \rho_2 = \dots = \rho_{10} = 0$

At significance level $\alpha = 0.05$, there is sufficient evidence to reject H_0 .

Serial correlation in errors: Error in one period is correlated with the errors in other periods.

Serial Correlation

- When fitting a regression model to time series data, it is common to find autocorrelation in the residuals.
- In this case, the estimated model violates the assumption of no autocorrelation in the errors (i.e., errors are independent over time).
 - Our forecasts may be inefficient — there is some information left over which should be accounted for in the model in order to obtain better forecasts.
 - The most common reason for serial correlation is that an important explanatory variable has been *omitted*.
- The forecasts from a model with autocorrelated errors are still unbiased, and so are not “wrong”, but they will usually have larger prediction intervals than they need to.



5. Multicollinearity

- In multiple-regression analysis, one of the assumptions that are made is that the independent variables are not highly linearly correlated with each other or with linear combinations of other independent variables.
- If this assumption is violated, a problem known as ***multicollinearity*** results.

	DPIPC	GDP	IR
DPIPC	1		
GDP	0.99	1	
IR	-0.65	-0.67	1

- If we look at these correlations between these variables, we can see a problem.
- Clearly, there is a very strong linear association between GDP and DPIPC.
- In this case, both of these variables are measuring essentially the same thing.


Example

5. Multicollinearity

- Note that no pair of independent variables has a correlation that is close to 1.0.
- The correlations of 1.00 along the main diagonal are to be expected.
 - They simply tell us that each variable is perfectly correlated with itself.

```
ncs %>%  
  GGally::ggpairs(columns = 3:6)
```

Excel: Data Analysis -> Correlation



Multiple Regressions: Things to Keep in Mind

- Independent variables should not only be statistically valid but also theoretically sound.
 - Correlation is not causation
- Monitor regularly the performance of a model. It may be the time to re-specify it.
- Statistics (e.g., adjusted R^2 , F and t) should be used as guide, not bible.
 - At times, you may find that even a model with one or two unsatisfactory statistics gives forecasts that are acceptable.
- **Principle of Parsimony:** If two models yield more or less the same result, the one with fewer variables should be preferred.

“Essentially, all models are wrong, but some are useful”

Your model has to be wrong...

... but that's o.k.
if it's illuminating!



George E.P. Box

Model selections

Useful predictors

Selecting predictors

- Dependent variable: Salary_K
- All possible models: How many?

	Years_in_Majors	Career_ERA	Innings_Pitched	Career_Wins	Career_Losses
1	1	0	0	0	0
2	1	0	1	0	0
3	1	1	1	1	1
			...		

Model Selection

- **Best subset selection** or “all possible subsets” regression
 - Recall from slide 01, suppose that there are k predictors, then the number of possible models is $2^k - 1$
- **Stepwise regression**
 - Backwards stepwise regression
 - Forward stepwise regression
 - Hybrid procedure
- Stepwise regression is a greedy heuristic algorithm
 - Not guaranteed to lead to the best possible model.



Demand Forecasting: Independent Variables (Predictors/Features)

- Internal/Endogeneous

- Price of product
- Media spending
- Number of new products to be introduced
- Number of products to be discontinued
- Number of stores
- FSI (Free Standing Insert) coupons distributed


- External/Exogeneous

- Competitive price
- New competitive product to be launched
- Gross Domestic Product (GDP)
- Consumer price index (CPI)
- Unemployment



Positive or negative relationship?

“Correlation does not imply causation”



<http://tylervigen.com/spurious-correlations>



Useful predictors for time series data

- Trend $y_t = \beta_0 + \beta_1 t + \epsilon_t$ where $t = 1, 2, \dots, T$

- Dummy variables for categorical predictor (later week)

- Trading days

- # of trading days in a month can vary considerably and can have a substantial effect on sales data.

- An alternative that allows for the effects of different days of the week has the following predictors

$$\left\{ \begin{array}{l} x_1 = \text{number of Mondays in month;} \\ x_2 = \text{number of Tuesdays in month;} \\ \vdots \\ x_7 = \text{number of Sundays in month.} \end{array} \right.$$

- Distributed lags (later week)

$$\begin{array}{l} x_1 = \text{advertising for previous month;} \\ x_2 = \text{advertising for two months previously;} \\ \vdots \\ x_m = \text{advertising for } m \text{ months previously.} \end{array}$$

- Fourier series (later week)

$$\begin{array}{l} x_{1,t} = \sin\left(\frac{2\pi t}{m}\right), x_{2,t} = \cos\left(\frac{2\pi t}{m}\right), x_{3,t} = \sin\left(\frac{4\pi t}{m}\right), \\ x_{4,t} = \cos\left(\frac{4\pi t}{m}\right), x_{5,t} = \sin\left(\frac{6\pi t}{m}\right), x_{6,t} = \cos\left(\frac{6\pi t}{m}\right), \end{array}$$



Selecting Predictors

1. Adjusted R^2 , $\bar{R}^2$

- Widely used; has been around longer than other measures.
- Maximizing $\bar{R}^2$ is equivalent to minimizing $\hat{\sigma}_e^2$

$$\hat{\sigma}_e^2 = \frac{1}{T - k - 1} \sum_{t=1}^T e_t^2$$

2. Cross-validation (CV)

- T = number of observations
- For large values of T , minimizing AIC is equivalent to minimizing CV value (MSE).

3. Information criterion: AIC, AICc, BIC

- Many statisticians like to use BIC. The textbook uses AICc.
- If the value of T is large enough, they will all lead to the same model.

AIC, AICc, BIC

- Akaike's Information Criterion (AIC)
 - The model with the minimum value of the AIC is often the best model for forecasting.
 - AIC penalizes terms more heavily than $\bar{R}^2$
 - Minimizing AIC is asymptotically equivalent to minimizing MSE via leave-one-out cross-validation.
- Corrected Akaike's Information Criterion (AICc)
 - For small values of T , the AIC tends to select too many predictors, and so AICc has been developed.
- Schwarz's Bayesian Information Criterion (BIC)
 - BIC penalizes the number of parameters more heavily than AIC.

T = number of observations

k = number of predictors in the model

SSE = sum of squared error

$$SSE = \sum_{t=1}^T e_t^2$$

$$AIC = T \log \left(\frac{SSE}{T} \right) + 2(k + 2);$$

$$AIC = -2 \log(L) + 2(k + 2)$$

where L is likelihood

$$AIC_c = AIC + \frac{2(k + 2)(k + 3)}{T - k - 3}$$

$$BIC = T \log \left(\frac{SSE}{T} \right) + (k + 2) \log(T)$$



Model selection

```
ncs %>%  
  mutate(URPR = UR*PR) %>%  
  model(  
    tslm0 = TSLM(NCS ~ DPIPC),  
    tslm1 = TSLM(NCS ~ DPIPC + UR + PR + UMICS),  
    tslm2 = TSLM(NCS ~ DPIPC + UR + PR + UMICS + URPR)  
  ) %>%  
  glance() %>%  
  select(.model, adj_r_squared, CV, AIC, AICc, BIC)
```

	.model	adj_r_squared	CV	AIC	AICc	BIC
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	tslm0	0.215	598194210.	1215.	1215.	1221.
2	tslm1	0.816	153330402.	1131.	1132.	1143.
3	tslm2	0.813	162755872.	1133.	1135.	1147.

Based on adjusted R^2 , we select _____

Based on MSE from cross validation, we select _____

Based on AICc, we select _____

Homework

Hw:

US Personal Consumption and Income

We would like to forecast changes in consumption (y) based on changes in income (x).

Obtain the regression equation

Perform diagnostic test

1. Do signs of coefficients make sense?
2. At significance level $\alpha = 0.05$, is the coefficient of income statistically significant?
3. Explanatory power
4. Residual analysis

```
#####
# Homework
#####
us_change
```

```
help(us_change)
```

```
> us_change
# A tibble: 198 x 6 [1q]
  Quarter Consumption Income Production Savings Unemployment
  <qtr>      <dbl>    <dbl>      <dbl>    <dbl>      <dbl>
1 1970 Q1      0.619    1.04      -2.45     5.30        0.9
2 1970 Q2      0.452    1.23      -0.551    7.79        0.5
3 1970 Q3      0.873    1.59      -0.359    7.40        0.5
```


us_change {fpp3}

R Documentation


Percentage changes in economic variables in the USA.

Description

us_change is a quarterly 'tsibble' containing percentage changes in quarterly personal consumption expenditure, personal disposable income, production, savings and the unemployment rate for the US, 1970 to 2016. Original \$ values were in chained 2012 US dollars.



```
us_change %>%
  pivot_longer(c(Consumption, Income),
               names_to = "var", values_to = "value") %>%
  ggplot(aes(x = Quarter, y = value)) +
  geom_line() +
  facet_grid(vars(var), scales = "free_y") +
  labs(title = "US consumption and personal income",
       y = "Quarterly % change")
```



```
us_change %>%
  ggplot(aes(x = Income, y = Consumption)) +
  geom_point() +
  labs(x = "Income",
       y = "Consumption",
       title = "Quarterly changes in US consumption and
personal income")
```

Appendix

Multiple Regression: Matrix Notation

Regression: Matrix notation

- Recall that multiple regression can be written as $y_t = \beta_0 + \beta_1 x_{1,t} + \dots + \beta_k x_{k,t} + \epsilon_t$
- Assume ϵ_t has mean zero and variance σ^2 .
- Let $\mathbf{y} = (y_1, \dots, y_n)^\top$, $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)^\top$, $\boldsymbol{\epsilon} = (\epsilon_1, \dots, \epsilon_n)^\top$.
 - T rows reflecting the number of observations
 - $k + 1$ columns reflecting the intercept which is represented by the column of ones plus the number of predictors

- The **design (model) matrix**

$$\mathbf{X} = \begin{bmatrix} 1 & | & & | & & | \\ \vdots & \mathbf{x}_1 & \dots & \mathbf{x}_i & \dots & \mathbf{x}_k \\ 1 & | & & | & & | \end{bmatrix} = \begin{bmatrix} 1 & x_{1,1} & x_{2,1} & \dots & x_{k,1} \\ 1 & x_{1,2} & x_{2,2} & \dots & x_{k,2} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & x_{1,T} & x_{2,T} & \dots & x_{k,T} \end{bmatrix}$$

- Matrix formulation: $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$

where $\boldsymbol{\epsilon}$ has mean $\mathbf{0}$ and variance $\sigma^2 \mathbf{I}$.

- Least squares estimation: Minimize $\boldsymbol{\epsilon}^\top \boldsymbol{\epsilon} = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$



Design (or Model) Matrix

Example: New Car Sales

- The model matrix

$$X = \begin{bmatrix} 1 & | & & | & & | \\ \vdots & x_1 & \dots & x_i & \dots & x_k \\ 1 & | & & | & & | \end{bmatrix}$$

File: 08-TS-Regression
sheet: matrix

NCS		ones	DPIPC	UR	PR	UMICS
152641		1	27840	5.7	4.75	93.13
166036		1	28134	5.83	4.75	94.1
175863		1	28168	5.73	4.75	87.27
151219		1	28363	5.87	4.45	83.83
152843		1	28584	5.87	4.25	79.97
173360		1	28958	6.13	4.24	89.27
178086		1	29539	6.13	4	89.3
154916		1	29708	5.83	4	91.97

```
dncs <- read_excel("09-TS-regression.xlsx", sheet = "NewCarSales")  
X <- model.matrix(NCS ~ DPIPC + UR + PR + UMICS, dncs)
```

	(Intercept)	DPIPC	UR	PR	UMICS
1	1	27840	5.70	4.75	93.13
2	1	28134	5.83	4.75	94.10
3	1	28168	5.73	4.75	87.27
4	1	28363	5.87	4.45	83.83
5	1	28584	5.87	4.25	79.97
6	1	28958	6.13	4.24	89.27

Multiple Regression: Matrix Notation

Least squares estimates

- Matrix formulation: $\mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \boldsymbol{\epsilon}$
- **Least squares estimation:** Minimize $\boldsymbol{\epsilon}^\top \boldsymbol{\epsilon} = (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X} \boldsymbol{\beta})$
- The least squares estimate $\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$
- This requires the inversion of the matrix $(\mathbf{X}^\top \mathbf{X})$
- If X is not of full column rank, then matrix $(\mathbf{X}^\top \mathbf{X})$ is singular, and the model cannot be estimated.
- The residual variance is estimated using

$$\hat{\sigma}_e^2 = \frac{1}{T - k - 1} (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}})' (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}}).$$

Multiple Regression: Matrix Notation

Least squares estimate

- Recall that multiple regression can be written as $y_t = \beta_0 + \beta_1 x_{1,t} + \dots + \beta_k x_{k,t} + \epsilon_t$
- Let $\mathbf{y} = (y_1, \dots, y_T)^\top$, $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)^\top$, $\boldsymbol{\epsilon} = (\epsilon_1, \dots, \epsilon_T)^\top$,
- Matrix formulation: $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$
- Least squares estimation:** Minimize $\boldsymbol{\epsilon}^\top \boldsymbol{\epsilon} = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$
- The least squares estimate $\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$



	A	B	C	D	E	F	G	H	I	J	K	L	M
1	NCS		ones	DPIPC	UR	PR	UMICS						
2	152641		1	27840	5.7	4.75	93.13			beta0	25428.26		
3	166036		1	28134	5.83	4.75	94.1			beta1	3.227462		
4	175863		1	28168	5.73	4.75	87.27			beta2	-7666.85		
5	151219		1	28363	5.87	4.45	83.83			beta3	-3219.5		
6	152843		1	28584	5.87	4.25	79.97			beta4	1122.429		
7	173360		1	28958	6.13	4.24	89.27						
8	178086		1	29539	6.13	4	89.3						
9	154916		1	29708	5.83	4	91.97						
10	159086		1	30094	5.7	4	98						

```
dncs <- read_excel("09-TS-regression.xlsx",
                    sheet = "NewCarSales")
X <- model.matrix(NCS ~ DPIPC + UR + PR + UMICS, dncs)
y <- data.matrix(dncs$NCS)
((solve(t(X) %*% X)) %*% t(X)) %*% y
```

```
      [,1]
(Intercept) 25428.255372
DPIPC        3.227462
UR          -7666.851578
PR          -3219.499125
UMICS        1122.429026
```

Multiple Regression: Matrix Notation

Forecasts and prediction intervals

Let $\mathbf{x}^*$ be a row vector containing the values of the predictors (in the same format as $\mathbf{X}$) for which we want to generate a forecast . Then the forecast is given by

$$\hat{y} = \mathbf{x}^* \hat{\boldsymbol{\beta}} = \mathbf{x}^* (\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Y}$$

and its estimated variance is given by

$$\hat{\sigma}_e^2 [1 + \mathbf{x}^* (\mathbf{X}' \mathbf{X})^{-1} (\mathbf{x}^*)'] .$$

A 95% prediction interval can be calculated (assuming normally distributed errors) as

$$\hat{y} \pm 1.96 \hat{\sigma}_e \sqrt{1 + \mathbf{x}^* (\mathbf{X}' \mathbf{X})^{-1} (\mathbf{x}^*)'} .$$

This takes into account the uncertainty due to the error term ε and the uncertainty in the coefficient estimates. However, it ignores any errors in $\mathbf{x}^*$. Thus, if the future values of the predictors are uncertain, then the prediction interval calculated using this expression will be too narrow.